

Example 6.2 - Ordered-categorical indicator variables

We are going to look at a dataset containing responses to five items of the LSAT Figure Classification test:

```
library("psych")
head(lsat6)
```

```
##      Q1 Q2 Q3 Q4 Q5
## [1,]  0  0  0  0  0
## [2,]  0  0  0  0  0
## [3,]  0  0  0  0  0
## [4,]  0  0  0  0  1
## [5,]  0  0  0  0  1
## [6,]  0  0  0  0  1
```

Part I: CFA approach (least-squares estimation)

First, let's look at the tetrachoric correlations:

```
tet <- tetrachoric(lsat6)
```

```
tet
```

```
## Call: tetrachoric(x = lsat6)
## tetrachoric correlation
##      Q1   Q2   Q3   Q4   Q5
## Q1 1.00
## Q2 0.17 1.00
## Q3 0.23 0.19 1.00
## Q4 0.11 0.11 0.19 1.00
## Q5 0.07 0.17 0.11 0.20 1.00
##
## with tau of
##      Q1   Q2   Q3   Q4   Q5
## -1.43 -0.55 -0.13 -0.72 -1.13
```

```
round(cor(lsat6), digits = 3)
```

```
##      Q1   Q2   Q3   Q4   Q5
## Q1 1.000 0.074 0.099 0.044 0.024
## Q2 0.074 1.000 0.115 0.062 0.086
## Q3 0.099 0.115 1.000 0.109 0.053
## Q4 0.044 0.062 0.109 1.000 0.099
## Q5 0.024 0.086 0.053 0.099 1.000
```

We see that the tetrachoric correlations, which account for the binary nature of the items, are higher than the Pearson correlations (calculated with `cor()`). This is what we generally see: for binary items, the Pearson correlation underestimates the strength of associations.

The tetrachoric correlation matrix also provides us with item thresholds. Lower values indicate easier items. We can also inspect the item means (i.e., proportion of respondents who had the items correct):

```
sort(colMeans(lsat6), decreasing = TRUE)
```

```
##      Q1      Q5      Q4      Q2      Q3
## 0.924 0.870 0.763 0.709 0.553
```

```
sort(tet$tau)
```

```
##           Q1           Q5           Q4           Q2           Q3
## -1.4325027 -1.1263911 -0.7159860 -0.5504657 -0.1332445
```

Both thresholds and item means indicate the same ordering of items in terms of difficulty. The most difficult item is Q3, the easiest item is Q1.

Let's perform a CFA using lavaan:

```
library("lavaan")
model.CFA <- '
  Theta =~ Q1 + Q2 + Q3 + Q4 + Q5
'
fit.CFA <- cfa(model.CFA, data = data.frame(lsat6), ordered = paste0("Q", 1:5))
paste0("Q", 1:5)
```

```
## [1] "Q1" "Q2" "Q3" "Q4" "Q5"
```

```
summary(fit.CFA, standardized = TRUE, fit.measures = TRUE)
```

```
## lavaan 0.6-18 ended normally after 29 iterations
##
##      Estimator                      DWLS
##      Optimization method            NLMINB
##      Number of model parameters      10
##
##      Number of observations          1000
##
## Model Test User Model:
##
##      Standard      Scaled
##      Test Statistic    4.051    4.740
##      Degrees of freedom      5      5
##      P-value (Chi-square)    0.542    0.448
##      Scaling correction factor      0.867
##      Shift parameter          0.070
##      simple second-order correction
##
## Model Test Baseline Model:
##
```

```

## Test statistic                    67.171    65.104
## Degrees of freedom                10        10
## P-value                          0.000      0.000
## Scaling correction factor         1.038
##
## User Model versus Baseline Model:
##
## Comparative Fit Index (CFI)       1.000      1.000
## Tucker-Lewis Index (TLI)         1.033      1.009
##
## Robust Comparative Fit Index (CFI) 1.000
## Robust Tucker-Lewis Index (TLI)   1.032
##
## Root Mean Square Error of Approximation:
##
## RMSEA                            0.000      0.000
## 90 Percent confidence interval - lower 0.000      0.000
## 90 Percent confidence interval - upper 0.039      0.043
## P-value H_0: RMSEA <= 0.050        0.989      0.982
## P-value H_0: RMSEA >= 0.080        0.000      0.000
##
## Robust RMSEA                      0.000
## 90 Percent confidence interval - lower 0.000
## 90 Percent confidence interval - upper 0.101
## P-value H_0: Robust RMSEA <= 0.050 0.690
## P-value H_0: Robust RMSEA >= 0.080 0.126
##
## Standardized Root Mean Square Residual:
##
## SRMR                             0.036      0.036
##
## Parameter Estimates:
##
## Parameterization                  Delta
## Standard errors                   Robust.sem
## Information                       Expected
## Information saturated (h1) model Unstructured
##
## Latent Variables:
##
## Estimate Std.Err z-value P(>|z|) Std.lv Std.all
## Theta =~
## Q1        1.000
## Q2        1.020    0.358    2.846    0.004    0.397    0.397
## Q3        1.210    0.447    2.709    0.007    0.471    0.471
## Q4        0.968    0.352    2.751    0.006    0.377    0.377
## Q5        0.879    0.352    2.499    0.012    0.342    0.342
##
## Thresholds:
##
## Estimate Std.Err z-value P(>|z|) Std.lv Std.all
## Q1|t1     -1.433    0.059   -24.431    0.000   -1.433   -1.433
## Q2|t1     -0.550    0.042   -13.133    0.000   -0.550   -0.550
## Q3|t1     -0.133    0.040    -3.349    0.001   -0.133   -0.133
## Q4|t1     -0.716    0.044   -16.430    0.000   -0.716   -0.716
## Q5|t1     -1.126    0.050   -22.395    0.000   -1.126   -1.126

```

```
##
## Variances:
##           Estimate Std.Err z-value P(>|z|) Std.lv Std.all
##      .Q1           0.848           0.848   0.848
##      .Q2           0.842           0.842   0.842
##      .Q3           0.778           0.778   0.778
##      .Q4           0.858           0.858   0.858
##      .Q5           0.883           0.883   0.883
##      Theta        0.152    0.087   1.743   0.081   1.000   1.000
```

We see that because we declared the items as ordered-categorical, the DWLS (diagonally weighted least squares) estimator was used. This provides standard and robust (robust against deviations from normality) fit indices, by default.

The model fits very well according to all fit measures. All loadings are substantial and significant. The most difficult item is Q3, easiest item is Q1. Also, Q3 is the strongest indicator of the latent trait, Q5 the weakest indicator.

Part II: IRT approach (maximum likelihood estimation)

Using R package `ltm`, we can perform a similar analysis, but not using ML estimation.

```
library("ltm")
lsat.IRT <- ltm(lsats6 ~ z1)
summary(lsat.IRT)
```

```
##
## Call:
## ltm(formula = lsats6 ~ z1)
##
## Model Summary:
##      log.Lik      AIC      BIC
## -2466.653 4953.307 5002.384
##
## Coefficients:
##           value std.err z.vals
## Dffc1t.Q1 -3.3597  0.8669 -3.8754
## Dffc1t.Q2 -1.3696  0.3073 -4.4565
## Dffc1t.Q3 -0.2799  0.0997 -2.8083
## Dffc1t.Q4 -1.8659  0.4341 -4.2982
## Dffc1t.Q5 -3.1236  0.8700 -3.5904
## Dscrmn.Q1  0.8254  0.2581  3.1983
## Dscrmn.Q2  0.7229  0.1867  3.8721
## Dscrmn.Q3  0.8905  0.2326  3.8281
## Dscrmn.Q4  0.6886  0.1852  3.7186
## Dscrmn.Q5  0.6575  0.2100  3.1306
##
## Integration:
## method: Gauss-Hermite
## quadrature points: 21
##
## Optimization:
```

```
## Convergence: 0
## max(|grad|): 0.024
## quasi-Newton: BFGS
```

The difficulty parameters reveal a similar ordering of item difficulty as the thresholds we estimated earlier. The discrimination parameters reveal a similar (but not completely identical) ordering of indicator strength as the loadings we estimated earlier.

Part III: Comparing the fit of the Rasch (1PL) and 2PL model

In the Rasch model, the probability of a correct answer is a function of the subject's ability and the item's difficulty:

$$p(Y = 1|\theta_j, \beta_i) = \frac{e^{\theta_j - \beta_i}}{1 + e^{\theta_j - \beta_i}}$$

where θ_j is the ability of person j , and β_i is the difficulty of item i .

In the 2PL model, the probability of a correct answer is additionally determined by the item's discriminatory power:

$$p(Y = 1|\theta_j, \beta_i, \alpha_i) = \frac{e^{\alpha_i(\theta_j - \beta_i)}}{1 + e^{\alpha_i(\theta_j - \beta_i)}}$$

where α_i is the discrimination parameter of item i .

We can empirically decide between the Rasch and 2PL model, by fitting both models to the data, and testing the difference in model fit.

We can do that using DWLS estimation in **lavaan**:

```
model.rasch <- '
  Theta =~ lambda*Q1 + lambda*Q2 + lambda*Q3 + lambda*Q4 + lambda*Q5
'
```

Note that I pre-multiplied all items with **lambda**. As a result, every item's loading will receive the same label, and all loadings will have the same estimated value. In effect, this applies an equality restriction on the item loadings.

We fit the model to the data and inspect the results:

```
fit.rasch <- cfa(model.rasch, data = data.frame(lsat6), ordered = paste0("Q", 1:5))
summary(fit.rasch, standardized = TRUE, fit.measures = TRUE)
```

```
## lavaan 0.6-18 ended normally after 2 iterations
##
##      Estimator              DWLS
##      Optimization method      NLMINB
##      Number of model parameters              6
##
##      Number of observations              1000
##
## Model Test User Model:
```

```

##                               Standard      Scaled
## Test Statistic                4.943        5.350
## Degrees of freedom              9           9
## P-value (Chi-square)           0.839        0.803
## Scaling correction factor       0.961
## Shift parameter                0.209
##   simple second-order correction
##
## Model Test Baseline Model:
##
## Test statistic                67.171        65.104
## Degrees of freedom             10           10
## P-value                       0.000        0.000
## Scaling correction factor       1.038
##
## User Model versus Baseline Model:
##
## Comparative Fit Index (CFI)     1.000        1.000
## Tucker-Lewis Index (TLI)       1.079        1.074
##
## Robust Comparative Fit Index (CFI) 1.000
## Robust Tucker-Lewis Index (TLI)  1.097
##
## Root Mean Square Error of Approximation:
##
## RMSEA                        0.000        0.000
## 90 Percent confidence interval - lower 0.000        0.000
## 90 Percent confidence interval - upper 0.021        0.023
## P-value H_0: RMSEA <= 0.050      1.000        1.000
## P-value H_0: RMSEA >= 0.080      0.000        0.000
##
## Robust RMSEA                  0.000
## 90 Percent confidence interval - lower 0.000
## 90 Percent confidence interval - upper 0.060
## P-value H_0: Robust RMSEA <= 0.050 0.915
## P-value H_0: Robust RMSEA >= 0.080 0.012
##
## Standardized Root Mean Square Residual:
##
## SRMR                        0.041        0.041
##
## Parameter Estimates:
##
## Parameterization              Delta
## Standard errors               Robust.sem
## Information                   Expected
## Information saturated (h1) model Unstructured
##
## Latent Variables:
##           Estimate Std.Err z-value P(>|z|) Std.lv Std.all
## Theta =~
## Q1      (lmbd)      1.000          0.400   0.400
## Q2      (lmbd)      1.000          0.400   0.400
## Q3      (lmbd)      1.000          0.400   0.400

```

```
##      Q4      (lmbd)    1.000                0.400    0.400
##      Q5      (lmbd)    1.000                0.400    0.400
##
## Thresholds:
##              Estimate Std.Err z-value P(>|z|) Std.lv Std.all
##      Q1|t1        -1.433   0.059  -24.431   0.000   -1.433   -1.433
##      Q2|t1        -0.550   0.042  -13.133   0.000   -0.550   -0.550
##      Q3|t1        -0.133   0.040   -3.349   0.001   -0.133   -0.133
##      Q4|t1        -0.716   0.044  -16.430   0.000   -0.716   -0.716
##      Q5|t1        -1.126   0.050  -22.395   0.000   -1.126   -1.126
##
## Variances:
##              Estimate Std.Err z-value P(>|z|) Std.lv Std.all
##      .Q1              0.840                0.840    0.840
##      .Q2              0.840                0.840    0.840
##      .Q3              0.840                0.840    0.840
##      .Q4              0.840                0.840    0.840
##      .Q5              0.840                0.840    0.840
##      Theta           0.160    0.025    6.341   0.000    1.000    1.000
```

We see good model fit according to all indices. Note that we have more degrees of freedom, because we estimated less parameters than in the previous model (Rasch model estimates 1 loading, the earlier model estimated 5 separate loadings for the items). We see that the standardized loadings are substantial and significant. The latent variable (`theta`) has significant variance. The ordering of item difficulties remained the same.

So should we prefer the more parsimonious Rasch model, or the more complex 2PL model? Although this is in large part a matter of personal preference (parsimony vs. complexity), we can also decide on statistical grounds, by comparing the fit indices:

```
fitinds <- c("chisq.scaled", "df", "pvalue.scaled", "cfi.scaled",
            "rmsea.scaled", "srmr")
fitMeasures(fit.CFA, fitinds)
```

```
##  chisq.scaled      df pvalue.scaled    cfi.scaled rmsea.scaled
##      4.740        5.000         0.448         1.000         0.000
##      srmr
##      0.036
```

```
fitMeasures(fit.rasch, fitinds)
```

```
##  chisq.scaled      df pvalue.scaled    cfi.scaled rmsea.scaled
##      5.350        9.000         0.803         1.000         0.000
##      srmr
##      0.041
```

Both models show excellent fit. Although χ^2 and SRMR indicate closer fit to the data for the 2PL model, the *df* indicate that the 1PL model is more parsimonious.

We can also statistically test the difference in model fit using a likelihood-ratio test:

```
lavTestLRT(fit.rasch, fit.CFA)
```

```
##
## Scaled Chi-Squared Difference Test (method = "satorra.2000")
##
## lavaan->lavTestLRT():
##   lavaan NOTE: The "Chisq" column contains standard test statistics, not the
##   robust test that should be reported per model. A robust difference test is
##   a function of two standard (not robust) statistics.
##           Df AIC BIC   Chisq Chisq diff Df diff Pr(>Chisq)
## fit.CFA    5      4.0511
## fit.rasch  9      4.9433    0.8764      4    0.9279
```

The likelihood ratio test indicates no significant difference in model fit between the 1- and 2PL model. In that case, we prefer the more parsimonious Rasch (1PL) model.

We could do the same comparison for the ML-estimated models:

```
lsat.IRT.rasch <- rasch(lsat6)
summary(lsat.IRT.rasch)
```

```
##
## Call:
## rasch(data = lsat6)
##
## Model Summary:
##      log.Lik      AIC      BIC
## -2466.938 4945.875 4975.322
##
## Coefficients:
##           value std.err   z.vals
## Dffclt.Q1 -3.6153  0.3266 -11.0680
## Dffclt.Q2 -1.3224  0.1422  -9.3009
## Dffclt.Q3 -0.3176  0.0977  -3.2518
## Dffclt.Q4 -1.7301  0.1691 -10.2290
## Dffclt.Q5 -2.7802  0.2510 -11.0743
## Dscrmn     0.7551  0.0694  10.8757
##
## Integration:
## method: Gauss-Hermite
## quadrature points: 21
##
## Optimization:
## Convergence: 0
## max(|grad|): 2.9e-05
## quasi-Newton: BFGS
```

```
anova(lsat.IRT.rasch, lsat.IRT)
```

```
##
## Likelihood Ratio Table
##           AIC      BIC log.Lik  LRT df p.value
## lsat.IRT.rasch 4945.88 4975.32 -2466.94
## lsat.IRT      4953.31 5002.38 -2466.65 0.57 4    0.967
```


Note that here we can compare models using information criteria (AIC, BIC). These information criteria are only defined for ML estimation, not for (DW)LS estimation (where we can still test difference using the the likelihood ratio test, or CFI, RMSEA, etc.). According to AIC and BIC, we should prefer the Rasch model. Furthermore, the likelihood ratio test does not indicate a difference in fit between the 1PL and 2PL model.

Part IV: Analysis of ordered categorical items with > 2 categories

For ordered items with > 2 ordered response categories, the code is the same. Just make sure you declare the items as ordered in applying the `cfa()` function. Automatically, a threshold for the number of categories - 1 is estimated. Reverse coding is not even necessary (items that should be reverse coded just get a negative loading, but you have to make sure that all categories within an item are ordered in the same direction).

With ordered-categorical items with > 2 categories, you can also compare the fit of a model in which all loadings are restricted to equality (i.e., the PCM or partial credit model) with a model in which all loadings are freely estimated (i.e., the GRM or graded response model). In `lavaan`'s `cfa()` function, you would do this by pre-multiplying the indicators of the latent trait by the same label, as we did above in `model.rasch`.

If you want to fit the GRM and PCM using ML estimation, you can use function `grm()` from package `ltm`. To fit the GRM model, use function `grm()` with and specify `constrained = FALSE`. To fit the PCM, use function `gram()` and specify `constrained = TRUE`.